

CS 7800 Information Retrieval

Simple Search Engine

Abhishek Singh Thakur
(U00933835)

Nikhil Chavan
(U00969241)

Code Snippets

Processing the text in util.py: Text would be processed in order to achieve tokenization, lower case, removal of stop words and stemming

```
import nltk
from nltk.stem.porter import *
import re
import norvig_spell
from norvig_spell import *

def tokenizer(text):
    text = re.sub("[^a-zA-Z]+", " ", text)
    tokens = nltk.tokenize.word_tokenize(text)
    return tokens

# ''' indexing a docuemnt, using the simple SPIMI algorithm, but no need to store
# blocks due to the small collection we are handling. Using save/load the whole in
# dex instead'''
# 1. Convert to lower cases,
# 2. Remove stopwords,
# 3. Stemming
def preprocessing_txt(text, query):
    tokens = text.split()
    # tokens = tokenizer(text)
    stemmer = PorterStemmer()
    processedText = ""
    removedStopwords = []
    stemmedWords = []
    if query == 'true':
        for i in range(len(tokens)):
            tokens[i] = correction(tokens[i])
    for token in tokens:
        token = token.lower()
        if token not in open('stopwords').read():
            processedText += stemmer.stem(token)
            processedText += " "
```

```

        if query == 'test':
            stemmedWords.append(stemmer.stem(token))
        elif query == 'test':
            removedStopwords.append(token)
    if query == 'test':
        print('Removed stop words are:' + str(removedStopwords))
        print('Stemmed words are:      ' + str(stemmedWords))
    return processedText

def isStopWord(word):
    n = 0
    for i in word.split():
        if i not in open('stopwords').read():
            n += 1
        if n == len(word.split()) - 1:
            print('No stop words founds')

```

Inverted Index creation Code implementation in -- index.py

```

# Indexing the Cranfield dataset and save the index to a file
# The index is saved to index_file and TFIDF scores are stored in tfidf scores for the terms in the documents from cran.all
def indexingCranfield(cranfile, savefile, test):
    # command line usage: "python index.py cran.all index_file"
    file = CranFile(cranfile)
    invertedIndex = InvertedIndex()
    temp = {}
    termFrequency = {}
    n = 0
    totalDocs = len(file.docs)
    print(str(len(file.docs)) + ' documents are present in the dataset.')
    for i in file.docs:
        if n == totalDocs: break
        else:
            # For calculating the Term frequency according to the documents
            temp = calculateTF(preprocessing_txt(i.body, ''))
            termFrequency[i.docID] = temp
            # For creating the index_file
            invertedIndex.indexDoc(i)
            n += 1
    invertedIndex.save(savefile, termFrequency, totalDocs, test)
    print('Indexing file creation is done')

class InvertedIndex:
    inverted_index = {}

```

```

inverted_index_temp = {}
def __init__(self):
    self.items = {} # list of IndexItems
    self.nDocs = 0 # the number of indexed documents

def indexDoc(self, doc): # indexing a Document object using functions in util
.py
    processedText = preprocessing_txt(doc.body, '')
    for word in processedText.split():
        position = wordPositions(word, processedText)
        if word in self.items.keys():
            if not int(doc.docID) in list(self.items[word].posting.keys()):
                self.items[word].add(int(doc.docID), position)
        else:
            index_item = IndexItem(word)
            index_item.add(int(doc.docID), position)
            self.items[word] = index_item
    for i in self.items:
        dictionary = {}
        x = 0
        for j in self.items[i].posting:
            listposition = []
            for k in self.items[i].posting[j]:
                listposition.append(k)
            dictionary[j] = listposition
        self.inverted_index_temp[i] = dictionary
    self.inverted_index = self.inverted_index_temp

# To find the normalized term location in the processed query
def wordPositions(word, processedText):
    pos = 1
    positions = []
    for data in processedText.split():
        if word == data:
            positions.append(pos)
        pos += 1
    return positions

```

Calculating TF-IDF for the document set cran.all which is implemented in index.py

```

def calculateIDF(doc, indexData, totalDocs):
    for i in doc.keys():
        if i in indexData.keys():

```

```

        count = len(indexData[i].keys())
        doc[i] = doc[i] * float( 1 + math.log10(totalDocs/count))
    return doc

# For calculating the normalized Term Frequency of the terms in the query
def calculateTF(text):
    textList = text.split()
    count = 0
    textTF = {}
    for i in range(len(textList)):
        if textList[i] in textTF.keys():
            textTF[textList[i]] += float(1/len(textList))
        else:
            textTF[textList[i]] = float(1/len(textList))
    return textTF

```

Saving Inverted Index in index_file and TF-IDF scores in the tfidf file for future use while calculating cosine similarity and relevant documents when using Boolean and Vector model for a sample query.

```

# Calculating and saving to disk TFIDF for the documents in the cran.all and crea
ting, serialize/deserialize the index indexing for the terms from all docs with t
heir positions
def save(self, savefile, termFrequency, totalDocs, test):
    for i in termFrequency:
        termFrequency[i] = calculateIDF(termFrequency[i], self.inverted_index
, totalDocs)
    # Storing tfidf scores for calculating Cosine similarity
    if test == 'true':
        with open('test_tfidf','w') as data:
            json.dump(termFrequency, data)
        with open('test_index','w') as fp:
            json.dump(self.inverted_index, fp)

    else:
        with open('tfidf','w') as data:
            json.dump(termFrequency, data)
        with open(savefile,'w') as fp:
            json.dump(self.inverted_index, fp)

```

Selecting of the Modes based on the command line input:

```
if mode == 0:
    booleanQuery(data, query, '')
elif mode == 1:
    vectorQuery(data, query, tfidf, '', 1400)
else:
    startBool = time.process_time()
    for i in sampleQueries:
        booleanQuery(data, i, 'time')
    print('Boolean Time      ' + str(time.process_time() - startBool))
    startVector = time.process_time()
    for i in sampleQueries:
        vectorQuery(data, i, tfidf, 'time', 1400)
    print('Vector Time      ' + str(time.process_time() - startVector))
```

Boolean Model Code Implementation: **query.py**

```
# Boolean query processing; note that a query like "A B C" is transformed to "A A
ND B AND C" for retrieving posting lists and merge them.
# Return a list of docIDs which matches with the query i.e. relevant documents fr
om cran.all
def booleanQuery(self, query, eval):
    processedText = preprocessing_txt(query, 'true')
    if eval == 'test':
        print('For making sure queries are converted to terms and the length
of the termslist is      ' + str(len(processedText.split())))
    queryDocs = []
    temp = []
    ndcgBoolean = []
    n = 0
    for i in processedText.split():
        if i in self:
            queryDocs.append(list(self[i].keys()))
    for i in range(len(queryDocs)):
        if i == 0:
            temp = queryDocs[i]
        else:
            temp = commonDocs(queryDocs[i], temp)
    if eval == 'boolean':
        if len(temp) == 0:
            noDocs = [0,0,0,0,0]
            return noDocs
        else:
            for i in range(5):
```

```

        if i < len(temp):
            ndcgBoolean.append(temp[i])
        else:
            ndcgBoolean.append(0)
        return ndcgBoolean
    print(ndcgBoolean)
elif eval == 'time':
    return
elif eval == 'test':
    if len(temp) > 0:
        print("The actual document which contains the query are --
" + str(temp)[1:-1])
    else:
        if len(temp) == 0:
            print('No Matching Results found')
        else:
            print("Boolean Model -
- " + str(len(temp)) + " documents which contains the query are " + str(temp)[1:-
1])
        return

# To find the common documents between two documents while calculating results for boolean model
def commonDocs(a, b):
    docs=[]
    for i in range(len(a)):
        for j in range(len(b)):
            if a[i] == b[j]:
                docs.append(a[i])
                break
    return docs

```

Vector Model implementation: **query.py**

```

# For vectorQuery, the program will output the top 3 most similar documents for a
query using the cosine similarity
# Returns top k pairs of (docID, similarity), ranked by their cosine similarity with
the query in the descending order
def vectorQuery(self, query, tfidf, eval, totalDocs):
    queryTerms = preprocessing_txt(query, 'true')
    if eval == 'test':
        print('For making sure queries in Vector Model are converted to terms and
the length of the termslist is' + str(len(queryTerms.split())))
    tffidfQuery, tffidfIndex, tfidfDocs, docTF = {}, {}, {}, {}
    # totalDocs = 1400

```

```

docSimilarity = []
tffidfQuery = tffidfQuery(queryTerms.split(), self, totalDocs)
tfidfIndex = tfidfIndex(queryTerms.split(), self, tfidf, totalDocs)
for doc in range(totalDocs):
    for i in queryTerms.split():
        tfidfDocs[i] = {}
        tfidfDocs[i] = tfidfIndex[i][str(doc+1)]
    temp = cosineSimilarity(tffidfQuery, tfidfDocs)
    docTF[temp] = {}
    docTF[temp] = str(doc + 1)
    docSimilarity.append(temp)
docSimilarity = sorted(docSimilarity, key = lambda x:float(x))

if eval == 'vector':
    ndcgVector = [0,0,0,0,0]
    cosine_value = [0,0,0,0,0]
    for i in range(5):
        if i < len(docSimilarity):
            ndcgVector[i] = docTF[docSimilarity[totalDocs-(i+1)]]
            cosine_value[i] = docSimilarity[totalDocs-(i+1)]
        else:
            ndcgVector[i] = 0
            cosine_value[i] = 0.0
    return ndcgVector+cosine_value
elif eval == 'time':
    return
else:
    print('Vector Model -
- Top 3 ranked Documents are ' + docTF[docSimilarity[totalDocs-1]] + ', ' +
      docTF[docSimilarity[totalDocs-2]] + ', ' + docTF[docSimilarity[totalDocs-
3]])
    print('          Their scores are          ' + str(docSimilarity[t
otalDocs-1]) + ', ' + str(docSimilarity[totalDocs-
2]) + ', ' + str(docSimilarity[totalDocs-3]))
# Calculating cosine similarity for ranking the documents against documents
def cosineSimilarity(querytfidf, doctfidf):
    modQuery = 0
    modDoc = 0
    dotProduct = 0
    for i in querytfidf.keys():
        modQuery += (querytfidf[i] * querytfidf[i])
        modDoc += (doctfidf[i] * doctfidf[i])
        dotProduct += (querytfidf[i] * doctfidf[i])
    if modQuery*modDoc == 0 :
        modDoc = 0.00000000000001

```

```

modQuery = 0.00000000000001
return float(dotProduct/(math.sqrt(modQuery) * math.sqrt(modDoc)))

```

Calculating TFIDF scores for the query sample, fetching the documents with the terms in the query sample:

```

# To calculate TDIDF scores for the query sample.
def tffIdfQuery(queryList, self, totalDocs):
    queryDict = {}
    for term in queryList:
        if term in queryDict.keys():
            if term in self:
                queryDict[term] += ((1/len(queryList)) * (1 + math.log10(totalDocs/len(self[term].keys()))))
            else:
                queryDict[term] += ((1/len(queryList)) * (1 + math.log10(1)))
        else:
            queryDict[term] = {}
            if term in self:
                queryDict[term] = ((1/len(queryList)) * (1 + math.log10(totalDocs/len(self[term].keys()))))
            else:
                queryDict[term] = ((1/len(queryList)) * (1 + math.log10(1)))
    return queryDict

# Fetching the documents with the TFIDF scores stored in tfidf file for every token in the query sample
def tfIdfIndex(queryList, self, tf, totalDocs):
    dict = {}
    for i in queryList:
        dict[i] = {}
        if i in self:
            # print(str(self[i].keys()))
            for k in range(totalDocs):
                if str(k+1) in self[i].keys():
                    dict[i][str(k+1)] = 0
                    dict[i][str(k+1)] = tf[str(k+1)][i]
                else:
                    dict[i][str(k+1)] = 0
        else:
            for m in range(totalDocs):
                dict[i][str(m+1)] = 0
    return dict

```


Mapping the Query samples from query.text with the qrels.text since docID's are not continues in the query.text for easy evaluation of Boolean and vector model results against the actual results. **batch_eval.py**

```
sampleQueries = []
f = open(sys.argv[3])
file = loadCranQry(sys.argv[2])
with open(sys.argv[1]) as fp:
    data = json.load(fp)
with open('tfidf') as tfidfFile:
    tfidf = json.load(tfidfFile)
queryIds = list(file.keys())
count = int(sys.argv[4])
querySamples = random.choices(queryIds, k=count)
querySamples = list(set(querySamples))
print(querySamples)
for i in querySamples:
    sampleQueries.append(file[i].text)
querydict = {}
mappingDict = {}
for i in f:
    j = i.split()
    if j[0] in querydict.keys():
        querydict[j[0]].append(j[1])
    else:
        querydict[j[0]] = []
        querydict[j[0]].append(j[1])
n = 1
for i in file.keys():
    mappingDict[i] = []
    mappingDict[i] = querydict[str(n)]
    n += 1
```

Calculating NDGC_5 scores and the Wilcoxon results for evaluating Boolean Model and Vector Model results with the actual results: **batch_eval.py**

```
import metrics
from cranqry import *
import random
import json
from query import *
from metrics import *
from scipy import stats
```

```

# Evaluating the NDCG scores for Boolean model and Vector model results against a
ctual data to check the success rate.
# Prints average NDCG scores for the boolean model, Vector model and p_value sco
re for sample of queries of size k
def eval():
    queryMatchesB = []
    queryMatchesV = []
    ndcg_scoreB = []
    ndcg_scoreV = []
    for i in range(len(sampleQueries)):
        queryMatchesB = booleanQuery(data, sampleQueries[i], 'boolean')
        ndcg_scoreB.append(calculate_ndgc5(queryMatchesB, mappingDict[querySample
s[i]], 'boolean'))

        queryMatchesV = vectorQuery(data, sampleQueries[i], tfidf, 'vector', 1400
)
        ndcg_scoreV.append(calculate_ndgc5(queryMatchesV, mappingDict[querySample
s[i]], 'vector'))
    avgBoolean = sum(ndcg_scoreB) / len(ndcg_scoreB)
    avgVector = sum(ndcg_scoreV) / len(ndcg_scoreV)
    p_value = stats.wilcoxon(ndcg_scoreB, ndcg_scoreV)
    print('Average NDCG score for Boolean Model is ' + str(avgBoolean))
    print('Average NDCG score for Vector Model is ' + str(avgVector))
    print('P-Value calculated for Boolean Vector is ' + str(p_value))

# Calculating NDCG scores for first 5 results for boolean and vector model.
def calculate_ndgc5(docId, resultDocs, method):
    y_truth = [0,0,0,0,0]
    y_score = [0,0,0,0,0]
    for i in range(5):
        if docId[i] == 0:
            y_truth[i] = 0
            y_score[i] = 0
        else:
            if method == 'boolean':
                y_score[i] = 1
            else:
                y_score[i] = docId[5+i]
            if docId[i] in resultDocs:
                y_truth[i] = 1
            else:
                y_truth[i] = 0
    ndcgValue = ndcg_score(y_truth, y_score)
    return ndcgValue

```

